

# TERRAFORM LESSON

*Francisco Javier Moreno Diaz*  
[fmorenод@gmail.com](mailto:fmorenод@gmail.com)



# Agenda

- IaC
- Terraform
- Lifecycle (and its commands)
- Blocks: Providers, Resource, Data, etc.
- Setup
- Exercise

# Terms

Template or Source Code)

State file

Stack

# What is IaC ?

Infrastructure as Code (IaC) is the process of managing and provisioning cloud infrastructure with ~~machine-readable~~ definition files.

Think of it as executable documentation.

Provide a codified workflow to create infrastructure (including changes/updates – version and scale)

Integrate with application code workflows (Git, CI/CD tools)

Provide reusable (repeatable & auditable) modules for easy sharing and collaboration

Safely test changes using native-tools of IaC or 3rd party


Enforce security policy and organizational standards

To sum all, it offers speed, reliability, collaboration, rollback



INFRASTRUCTURE-  
AS-CODE

# What is Terraform?




HashiCorp  
**Terraform**

Declarative Open-source provisioning tool (IaC).  
It ships as a cross-platform single binary which is written in Go.  
Executable Documentation: Human and machine readable  
Easy to learn  
Test, share, re-use, automate  
Works on all major cloud providers



# Architecture




Provides a common language (HCL v2) to orchestrate cloud resources and/or other providers.

# Use case

HashiCorp


# Terraform





"Come on, make up your mind -  
or it's back to the Sinkhole of Nested XML."

# Workflow as example



# Terraform is ....

is a command line tool. It is the same for all platforms.

Terraform commands are either typed in manually or run automatically from a script.

Terraform has subcommands that perform different actions.

# Basic Terraform Commands

terraform version

terraform help

terraform init

terraform plan

terraform apply

terraform destroy

Type `terraform subcommand help` to view help on a particular subcommand.

# Terraform Code

```
resource aws_vpc "main" {  
  cidr_block      = "10.0.0.0/16"  
  instance_tenancy = "dedicated"  
}
```

Terraform code is based on the HCL2 toolkit. HCL stands for HashiCorp Configuration Language.

Terraform code, or simply `terraform` is a declarative language that is specifically designed for provisioning infrastructure on any cloud or platform.

# Terraform init

```
$ terraform init
```

Initializing provider plugins...

- Checking for available provider plugins...
- Downloading plugin for provider "aws" (hashicorp/aws) 2.35.0...

...

provider.aws: version = "~> 2.35"

Terraform has been successfully initialized!

Terraform fetches any required providers and modules and stores them in the .terraform directory. If you add, change or update your modules or providers you will need to run `terraform init` again.

# Terraform Plan

```
$ terraform plan
```

An execution plan has been generated and is shown below.

Terraform will perform the following actions:


```
# aws_vpc.main will be created
+ resource "aws_vpc" "main" {
    + arn                               = (known after apply)
    + cidr_block                         = "10.0.0.0/16"
    ...
    + instance_tenancy                   = "dedicated"
}
```

Preview your changes with `terraform plan` before you apply them.


# Terraform Dependency Graph

The terraform resource graph visually depicts dependencies between resources.

The region and prefix variables are required to create the resource group, which is in turn required to build the virtual network.



# Files



Depends on your extension, you reorganize the content on files, however it is a good advice to manage a similar structure

# Blocks

Providers

Variables

Resource

Data

Outputs

Parameters

Out of scope for this lesson:

Modules

Provisioner

Workspace

Remote backend

Lock

Loops or Conditions

Built-in Variables

# Terraform Provider Configuration

The terraform core program requires at least one provider to build anything. You can manually configure which version(s) of a provider you would like to use. If you leave this option out, Terraform will default to the latest available version of the provider.

```
provider "aws" {  
    version = "=2.35.0"  
}
```

- = (or no operator): exact version equality
- !=: version not equal
- \>, >=, <, <=: version comparison
- ~>: pessimistic constraint, constraining both the oldest and newest version allowed. ~> 0.9 is equivalent to >= 0.9, < 1.0, and ~> 0.8.4 is equivalent to >= 0.8.4, < 0.9

# Variables

## INPUT

Define variables: See priority next slide.

Types of variables:

- \* Simple: string, number, bool
- \* Constructed: list, set, map, object, tuple.
- \* Validation: regexp, null,

## OUTPUT

Attributes: Sensitive

```
main.tf > variable "region" > default
variable "tenancy_ocid" {
  description="el Identificador del Tenant"
  type=string
  default = ""
}

variables.env
set TF_VAR_tenancy_ocid=ocid1.tenancy/...
set TF_VAR_user_ocid=ocid1.user/...
set TF_VAR_fingerprint=1d:6d:f7:0e:...
set TF_VAR_compartment_ocid=ocid1.compartment/...

terrafrom.tfvars > region
private_key_path="C:\\\\Users\\\\framoren\\\\.ssh\\\\id_rsa"
ssh_public_key="C:\\\\Users\\\\framoren\\\\.ssh\\\\id_rsa.pub"
region="us-ashburn-1"
```

# Variables

Terraform variables are placed in a file called variables.tf. Variables can have default settings. If you omit the default, the user will be prompted to enter a value. Here we are declaring the variables that we intend to use.

## Priority (high to low)

1. Command line flag - run as a command line switch
2. Configuration file - set in your terraform. tfvars file
3. Environment variable - part of your shell environment
4. Default Config - default value in variables.tf
5. User manual entry - if not specified, prompt the user for entry

# Block Structure

```
provider "oci" {
    region      = var.region
    tenancy_ocid = var.tenancy_ocid
    user_ocid   = var.user_ocid
    fingerprint = var.fingerprint
    private_key_path = var.private_key_path
}

# See https://docs.oracle.com/iaas/images/
data "oci_core_images" "test_images" {
    compartment_id      = var.compartment_ocid
    operating_system    = "Oracle Linux"
    operating_system_version = "8"
    shape               = var.instance_shape
    sort_by             = "TIMECREATED"
    sort_order          = "DESC"
}

/* Network */
resource "oci_core_virtual_network" "test_vcn" {
    cidr_block      = "10.1.0.0/16"
    compartment_id = var.compartment_ocid
    display_name   = "testVCN"
    dns_label      = "testvcn"
}

output "Public_IP_LoadBalanceador" {
    value = "http://\${oci\_load\_balancer\_load\_balancer.free\_l
}
```

*variable  
provider  
data  
resource  
output*

## Module

Folder with code and has similar input and output

Terraform accesses TF files in the current folder only, does not enter to nested folders

# How to Provision an AWS Instance

Before we start, we'll need to gather some basic information including (but not limited to):

Instance Name, Operating System (Image), VM Size,  
Geographical Location (Region), Security Groups

```
resource aws_instance "web" {
    ami           = data.aws_ami.ubuntu.id
    instance_type = "t2.micro"

    tags = {
        Name = "HelloWorld"
    }
}
```

# Anatomy of a Resource

Every terraform resource is structured the same way.

**resource** = Top level keyword

**type** = Type of resource. Example: `aws_instance` .

**name** = Arbitrary name to refer to this resource. Used internally by terraform.  
This field cannot be a variable.

```
resource type "name" {
  parameter = "foo"
  parameter2 = "bar"
  list = ["one", "two", "three"]
}
```

# Data Sources

```
data "aws_ami" "ubuntu" {
  most_recent = true
  filter {
    name      = "name"
    values    = ["ubuntu/images/hvm-ssd/ubuntu-trusty-14.04-amd64-server-*"]
  }
  filter {
    name      = "virtualization-type"
    values    = ["hvm"]
  }
  owners    = ["099720109477"] # Canonical
}
```

Data sources are a way of querying a provider to return an existing resource, so that we can access its parameters for our own use.

# Terraform Dependency Mapping

Terraform can automatically keep track of dependencies for you. Look at the two resources below. Note the highlighted lines in the `aws_instance` resource. This is how we tell one resource to refer to another in terraform.

You can use [depends\\_on](#)

```
resource aws_key_pair "my-keypair" {
  key_name      = "my-keypair"
  public_key    = file(var.public_key)
}

resource "aws_instance" "web" {
  ami           = data.aws_ami.ubuntu.id
  instance_type = "t2.micro"
  key_name      = aws_key_pair.my-keypair.name
```

# Terraform State

Terraform is a **stateful** application. This means that it keeps track of everything you build inside of a state file. You may have noticed the `terraform.tfstate` and `terraform.tfstate.backup` files that appeared inside your working directory. The state file is Terraform's source of record for everything it knows about.

```
{  
  "terraform_version": "0.12.7",  
  "serial": 14,  
  "lineage": "452b4191-89f6-db17-a3b1-4470dcb00607",  
  "outputs": {  
    "catapp_url": {  
      "value": "http://go-hashicat-5c0265179ccda553.workshop.aws.hashidemos.io",  
      "type": "string"  
    },  
  }  
}
```

# Changing Existing Infrastructure

Whenever you run a plan or apply, Terraform **reconciles** three different data sources:

1. What you wrote in your code
2. The state file
3. What actually exists

Terraform does its best to add, delete, change, or replace existing resources based on what is in your \*.tf files. Here are the four different things that can happen to each resource during a plan/apply:


```
+  create
-  destroy
-/+ replace
~  update in-place
```

# Best Practices

- Keep configurations modular
- Use variables and workspaces for environments
- Never edit the state file manually
- Use `terraform fmt` & `terraform validate`
- Protect credentials (IAM roles, not keys)

# Setup

1. Install of terraform and validate
2. Install of AWS CLI and validation –optional-
3. Obtain AWS credentials for CLI
4. Folder with Terraform Code
5. Lifecycle: Init, Plan and Apply
6. Check working environment
7. Changes and Lifecycle
8. End of infrastructure: Destroy



# Official Tutorial

-   https://developer.hashicorp.com/terraform/tutorials/aws-get-started

  Terraform [Install](#) [Tutorials](#) [Documentation](#) [Sandbox](#) [Registry](#) [Try Cloud](#)

[Terraform Home](#)   [Developer](#) / [Terraform](#) / [Tutorials](#) / [AWS](#)



## Get Started - AWS

Create, manage, and destroy AWS infrastructure using Terraform. Step-by-step, command-line tutorials will walk you through the Terraform basics for the first time.

[Create an account](#) to track your progress.

[Start](#)  6 tutorials

 3min

**What is Infrastructure as Code with Terraform?**

Learn how infrastructure as code lets you safely build, change, and manage infrastructure. Try Terraform.

 7min

**Install Terraform**

Install Terraform on Mac, Linux, or Windows by downloading the binary or using a package manager (Homebrew or Chocolatey). Then

# References & Resources

- <https://developer.hashicorp.com/terraform/docs>
- <https://registry.terraform.io/providers/hashicorp/aws/latest/docs>
- <https://learn.hashicorp.com/collections/terraform/aws-get-started>
- <https://aws.amazon.com/free/>